


Step 3 - Multi-Evidence Object Tracking

Associate frame-local observations into persistent image-space object identities

Output: ID-consistent mask tracks



Missing cues are explicit; unavailable terms are omitted rather than treated as contradictory evidence.

Persistent IDs are assigned in Step 3, not Step 2.